

# Integrating IPFS ELK Stack for Resilient and Queryable Data Storage over Edge and Fog

Jeet Ahuja, Akshat Kumar, Yogesh Simmhan  
*Department of Computational and Data Sciences*  
*Indian Institute of Science*  
Bangalore, India

**Abstract**—Immense growth in the field of IoT and Edge has also brought with it an exponential increment in the amount of data that comes naturally because of the 1000s of sensors which export time-series data in real time. As such, it is often essential to query for such data for multiple purposes ranging from working with a subset of data to monitoring the health of the infrastructure. With the improvement of edge devices, we can now retain and store data locally on the edge devices itself rather than going through the hassle of moving every single bit of it to the central cloud. However, what such systems do lack is the availability of flexible querying over all such data which may be spread across 100s of such devices. This is where TorqueDB, developed by Dhruv Garg, et al. comes in handy. By default, TorqueDB leverages on the use of ElfStore, a distributed edge-local file store for file system and InfluxDB, a time-series database. For temporal querying on the data across the edge and fog devices. We intend to make the TorqueDB system more efficient by replacing ElfStore with IPFS (Inter-Planetary file system) and use Elasticsearch from ELK stack for querying and indexing over the data.

## I. INTRODUCTION

### A. Context

IoT domains span physical infrastructure such as Smart Cities, Smart Transportation and Industrial IoT, to consumer devices such as Smart watches and smart appliances. In IoT applications the data collected from an array of sensors is analysed and decisions are made within few seconds to control the system.

As of late, edge devices are being widely deployed, because they help gather and transmit observations from the sensors, and to enact control decisions. This has resulted in a lot of data, especially time-series data.

### B. Motivation

IoT applications are often deployed on an array of several sensors that collect huge amounts of time-series data that help monitor, analyse, and take appropriate action over the concerned events.

However, a boom in this field has seen an explosive growth in the amount of time-series data, generated especially by high-sampling frequency sensors. As such, the actions to be taken as an outcome of such data is generally a result of the analysis performed over a recent subset of the collected data. This is more flexible than Complex Event Processing (CEP)

and publish-subscribe systems that operate on data streaming data and limit possible queries.

Although, TSDBs like InfluxDB are popular choices for hosting such data from sensors and perform temporal queries on the data as these TSDBs reside on the cloud.

However, since the sensors producing this data and the application consuming this data tend to reside close to the edge devices, sending every query to the central cloud can incur significant latency which is an important factor to consider in cases where the Edge application requires sub-second query latency to respond to dynamic solutions. Also, there are the factors of unreliability, security and cost when moving data from edge to cloud and serving queries in the opposite direction.

Also, we lose a lot of compute and storage potential in terms of the resources available on the edge and fog devices themselves. This gives us the motivation to develop an efficient storage and querying system that resides right on the edge and fog devices.

### C. Contributions

This paper contributes by running an Elasticsearch over the IPFS storage. The data from sensors collected on edge devices is stored in the IPFS and then Elasticsearch is used as indexing and querying mechanism.

This combination can be used as a replacement of ElfStore. It is expected that IPFS with Elasticsearch performs better than ElfStore. Because Elasticsearch is mainly used for fast queries. And Elasticsearch does operations in a distributed manner. It performs Indexing and Querying Next, in *section-2*, we discuss the background of ElfStore and related works; in *section-3*, we will discuss the Background and Technical Approach, where IPFS, Elasticsearch, InfluxDB are discussed; in *section-4* we discuss the proposed architecture and the approaches that we tried.

## II. GAPS/RELATED WORK

The above discussed motivation is where TorqueDB developed by Dhruv Garg, et al. comes in handy. It leverages ElfStore distributed edge local storage and InfluxDB to offer

distributed time-series querying over the edge and fog devices

We intend to make changes to these leverages by replacing ElfStore with IPFS and use Elasticsearch (ELK Stack) as an indexing mechanism in an effort to improve the efficiency and usability of the database.

#### A. ElfStore

Elfstore is an edge-local federated store for streams of data blocks. i.e., it is stream-based, block-oriented distributed storage service over unreliable edges. Reliable fog devices are used as a super-peer overlay to monitor the edge resources.

Elfstore provides federated metadata indexing using Bloom filters. It maintains approximate global statistics about the reliability and storage capacity of edges.

In ElfStore there are two resources, edge and fog. The fog resources are reliable and they connect with each other through a WAN/MAN. ElfStore supports some of the service API like, CreateStream(), UpdateBlock(), PutBlock(), UpdateStreamMeta(), FindStream(), FindBlock(), etc. The data model for ElfStore is a stream of blocks.

ElfStore performs device management and search using a P2P model, where the fog acts as super-peer and edges act as peers among them. ElfStore will host the actual data and to guarantee reliability and balance storage utilization it performs replication.

### III. BACKGROUND AND TECHNICAL APPROACH

**IPFS:** The Interplanetary File System is a protocol and peer-to-peer network for storing and sharing data in a distributed file system.

A content addressed storage is employed by IPFS. i.e., for accessing data through IPFS, contents of data is being, instead of location based addressing

The mechanism is like BitTorrent. Instead of having a centrally located server, IPFS is built around a decentralized system of multiple host-operator nodes where each node holds a part of the total data, resulting in a resilient system of file storage and sharing. Any node in the network can serve the required file using the content as an address, and other peers can find and request that content from any host having that content using Distributed Hash Table textcomp

Some commands of IPFS include, *ipfs add x*, which adds file x to the filesystem and print the resulting hash of the file. *ipfs get <hash>* retrieves a file from the filesystem, where <hash> is the IPFS hash of the required file

In IPFS, data is stored using IPFS objects. The IPFS object format is,

```
type IPFSObject struct {
```

```
    links [] IPFSLink
    //array of links

    data [] byte
    // opaque content data
}
```

A single IPFS object can store 256KB of data. For storing files with size greater than 256KB, splitting of file is performed with each split being less than 256KB

In the case of splitting of file, an empty parent object links all the pieces of the file together using the link field.

commit objects are used to support versioning in IPFS. parent object is linked to the updated object and so all updates are tracked as well as the older versions are also maintained

For our project, we'll have an IPFS network consisting of multiple fogs and edge resources. Client will send a query to the fog co-ordinator and the fog co-ordinator node will publish the query on the network.

**Elasticsearch:** Elasticsearch is distributed, open-source search and analytics engine built using Apache Lucene and developed in Java. It is a scalable version of the Lucene search framework and has the added ability to horizontally scale Lucene indices. Elasticsearch allows us to store, search and analyse huge volumes of data quickly and efficiently by searching the text directly. In simpler terms, we can think of Elasticsearch as a server that receives JSON requests and gives back JSON data, all of it being leveraged with extensive REST APIs

It uses inverted index. i.e., a hashmap-like data structure that directs from a word to a document. Elastic search distributes the tasks of searching and indexing across all the nodes in the cluster

ElasticSearch ensures redundancy by implementing sharding and replication techniques.

Elastic Stack (ELK Stack) is a set of open-source tools for data ingestion, enrichment, storage, analysis and visualization. And elastic search is a component of ELK stack

So in our case ElasticSearch indexes based on the meta-data. And when a client queries, the searching is done based on the meta-data. First step will be to elastic search the index using the meta-data requested.

**LogStash** LogStash is a data processing pipeline used to process data on the fly and send it to the required destination. Also, it helps us obtain data in required format from a variety of different filetypes.

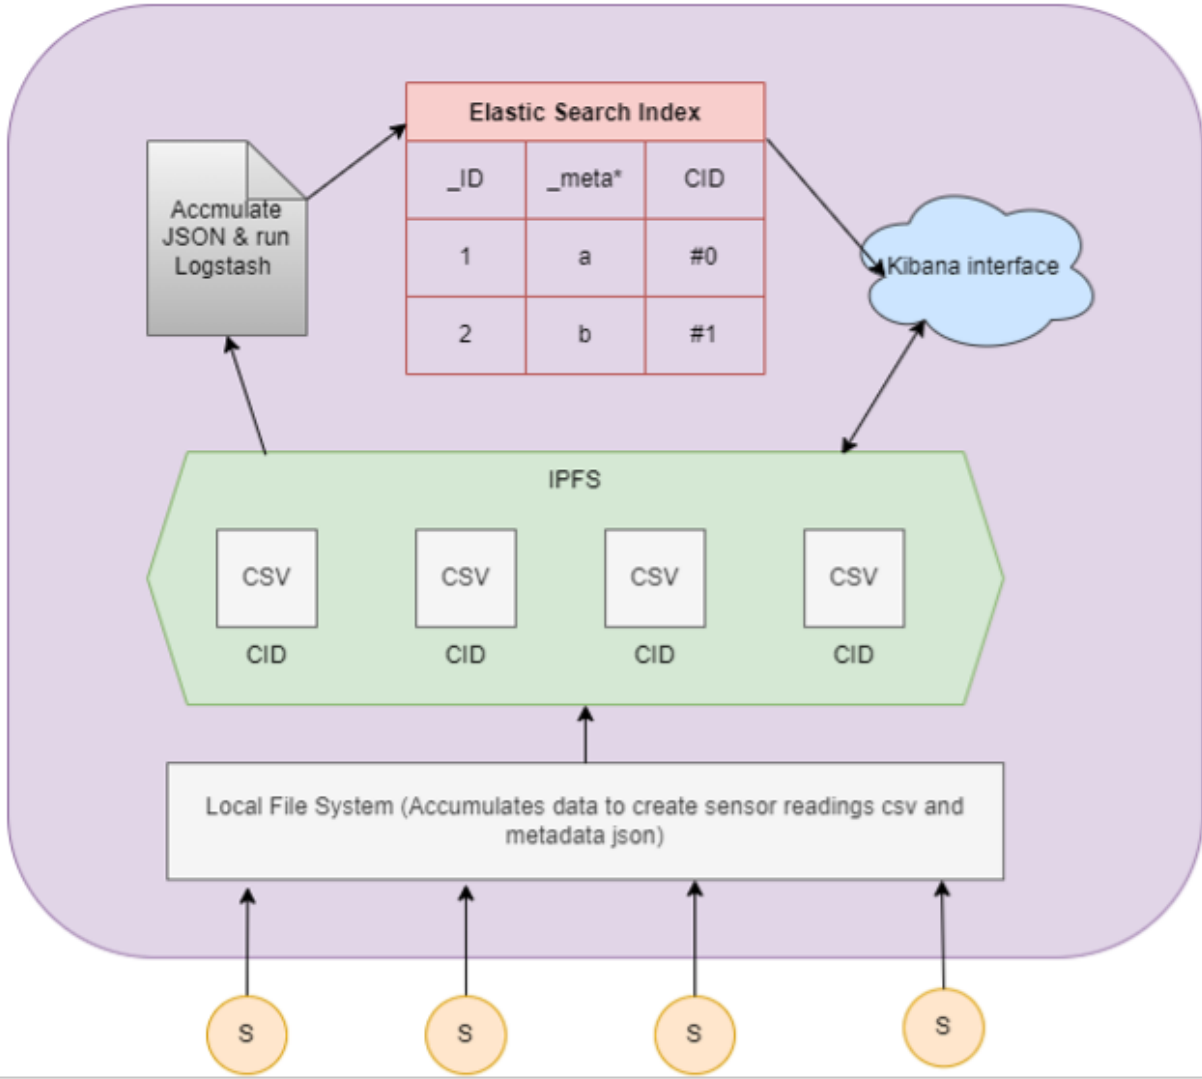


Fig. 1. High level Architecture

LogStash is a tool used to extract information even if the data is unstructured making use of variety of supported filter APIs and native Regex support.

**Kibana** A front-end application that fits on top of ELK Stack. It is used for visualization of the data indexed by ElasticSearch.

#### IV. PROBLEM DEFINITION

The focus of this project is to integrate IPFS and ELK Stack to get a powerful distributed data storage and data-queurying platform

First we need to integrate IPFS with ELK Stack on a per-node basis to get the basic functionality through.

For better query ability, the data stored should be indexed, so that a query can be executed in an efficient manner. So, for indexing ElasticSearch is used. ElasticSearch creates index based on metadata. First step is to identify using ElasticSearch the storage block id using the meta-data requested

#### V. PROPOSED ARCHITECTURE

1. Sensor data is accumulated by the edge device and a time-series file is generated with the required metadata
2. This time-series files are inserted into IPFS and after inserting, a unique CID (Content Identifieers) are generated. This CIDs are appended to metadata as they help retrieving the blocks using IPFS's API
3. Then aggolomerate all metadata into one single json and push it with logstash json filter into ElasticSearch, where it is

indexed

4. Perform basic queries on data via ELC CLI

## VI. OUR APPROACH

1. We first install and configure IPFS

2. Initialize IPFS using init command

3. Load the directory to be stored into IPFS

IPFS stores all the directory files by splitting each file into smaller chunks and returns a hash key called a content identifier

4. Retrieve the hash keys for each file and store in a txt file.

5. Use LogStash to form an index on metadata by writing a config file

6. Query the index using ElasticSearch

7. Visualize using Kibana

After this, integrate the IPFS to TorqueDB and modify the functions integrate the query mechanism, *readEdge()*, *find-BlocksAndLocationsWithQuery()* and *getMetadataByBlockid()*

## VII. FUTURE WORK

- Now that we have the basics down of how to integrate Elasticsearch with IPFS, we can now make necessary changes in the logstash bul push config file.
- The first change is to make id of Elasticsearch as the metadata of the file itself i.e. instead of indexing the metadata, we will index with the metadata.
- Append the value field to have “ipfs get” in front of the hash key and make it auto execute.
- This means that the user can search files in Elastic search using file metadata and in return, the search will initiate a retrieve file with that metadata and get the required files ready for further work in InfluxDB.

## VIII. ACKNOWLEDGEMENT

This project has been proposed for the coursework DS256: Scalable Systems for Data Sciences under the guidance of Dr. Yogesh Simmhan (Associate Professor, Department of Computational and Data Sciences, IISc Bangalore, India).

- <https://www.knowi.com/blog/what-is-elastic-search/>
- <https://en.wikipedia.org/wiki/InterPlanetaryFileSystem>
- Dhruv Garg, et al.: TorqueDB: Distributed Querying of Time-series Data from Edgelocal Storage, Euro-Par 2020, <https://doi.org/10.1007/978-3-030-57675-218>